

HW-10-PCA_Wes_Wilson

Packages

```
library(tidyverse)
library(ISLR2)
library(pls)
library(MASS)
library(leaps)
library(glmnet)
```

Data and Functions

```
data("Boston")
data("College")
split_train_and_test_data <- function(split_pct, data, seed){
  set.seed(seed)
  z <- sample(nrow(data), split_pct * nrow(data))
  train <- data[z,]
  test <- data[-z,]

  return(list(train = train, test = test))
}

calc_mse <- function(y_actual, y_predicted) {
  if (!is.vector(y_actual) || !is.vector(y_predicted)) {
    warning("Both y_actual and y_predicted should be vectors.")
    return(NULL) # Return NULL to indicate an issue
  }
  mse <- mean((y_actual - y_predicted)^2)

  return(mse)
```

```
}
```

College Applications Continued

This continues the College data problem from the previous h/w. (Chap. 6, # 9 cd, p.286-287)

Predict the number of applications received based on the other variables in the College data set.

- This data set is from our textbook. Access it with `library(ISLR2)`.

Now fit models using two additional regularization methods using the validation set based on 50% split where appropriate and `set.seed(123)` where appropriate:

```
model_data <- split_train_and_test_data(split_pct = .5,
                                         seed = 123,
                                         data = College)
```

(a) PCR model, with M, the number of principal components, chosen by cross-validation

```
pcr_reg <- pcr(Apps ~ ., data = model_data$train, validation = "CV")
summary(pcr_reg)
```

Data: X dimension: 388 17

Y dimension: 388 1

Fit method: svdpc

Number of components considered: 17

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	3258	3258	1482	1463	1342	1235	1076
adjCV	3258	3258	1481	1462	1339	1234	1073
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	1075	1082	1083	1065	1050	1053	1047
adjCV	1073	1080	1080	1062	1047	1049	1044
	14 comps	15 comps	16 comps	17 comps			
CV	1042	1039	1042	1023			
adjCV	1038	1035	1038	1018			

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	53.738	87.86	95.79	97.67	98.77	99.40	99.91	99.96
Apps	1.547	79.79	80.49	83.71	86.45	89.88	89.91	89.91
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	100.00	100.00	100.00	100.00	100.00	100.00	100.00	
Apps	89.92	90.36	90.65	90.69	90.85	91.05	91.19	
	16 comps	17 comps						
X	100.0	100.00						
Apps	91.2	91.58						

(b) PLS model, with M, the number of principal components, chosen by cross-validation.

```
pls_reg <- plsreg(Apps ~ ., data = model_data$train, validation = "CV")
summary(pls_reg)
```

Data: X dimension: 388 17
Y dimension: 388 1
Fit method: kernelpls
Number of components considered: 17

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	
CV	3258	1472	1430	1256	1121	1070	1066	
adjCV	3258	1466	1446	1256	1120	1068	1065	
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps	
CV	1072	1075	1056	1041	1034	1030	1028	
adjCV	1070	1073	1059	1038	1031	1026	1025	
	14 comps	15 comps	16 comps	17 comps				
CV	1029	1029	1026	1014				
adjCV	1025	1026	1023	1010				

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	35.02	86.48	93.11	97.29	98.47	99.40	99.77	99.96
Apps	80.22	81.34	85.77	88.84	89.82	89.91	89.93	89.95
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	99.99	100.00	100	100.00	100.00	100.0	100.0	
Apps	90.24	90.75	91	91.16	91.19	91.2	91.2	
	16 comps	17 comps						
X	100.00	100.00						

Apps 91.22 91.58

(c) Evaluate performance of each method in terms of **prediction MSE** and add the results into the summary table.

(d) Comment on the results obtained.

With the PCR model, around 6 components appears to be the optimal number. Once we move beyond 10 or so components, the changes in the variance explained by the x values and the rmse is nominal.

With the pls model, around 9 components looks to be an optimal number; though I suppose that is arguable depending on who you ask. I am going to create new models that have the optimal number of components and use that to predict and evaluate performance.

```
pcr_predictions <- predict(object = pcr_reg,
                           newdata = model_data$train, ncomp = 6) %>%
  as.vector()

plsr_predictions <- predict(object = pls_reg, newdata = model_data$train,
                           ncomp = 9) %>% as.vector()

# the pls functions already does this; but doing again to better understand.
calc_mse(y_actual = model_data$train$Apps,
        y_predicted = pcr_predictions)
```

```
[1] 1068315
```

```
calc_mse(y_actual = model_data$train$Apps,
        y_predicted = plsr_predictions)
```

```
[1] 1030691
```

Do on the test data now.

```
pcr_predictions <- pcr_predictions <- predict(object = pcr_reg,
                                              newdata = model_data$test, ncomp = 6) %>%
  as.vector()

plsr_predictions <- predict(object = pls_reg, newdata = model_data$test,
                           ncomp = 9) %>% as.vector()
```

```

pcr_rmse <- calc_mse(y_actual = model_data$test$Apps,
                    y_predicted = pcr_predictions) %>% sqrt()

plsr_rmse <- calc_mse(y_actual = model_data$test$Apps,
                    y_predicted = plsr_predictions) %>% sqrt()

```

- How accurately can we predict the number of college applications?

```

tibble(method = c("step_lse", "step_k_fold", "ridge", "lasso_lse", "lasso_min", "pcr"),
       results = c(sqrt(1327026), sqrt(1264962), sqrt(2092402),
                   sqrt(1528732), sqrt(1391186), pcr_rmse, plsr_rmse))

```

```

# A tibble: 7 x 2
  method      results
  <chr>      <dbl>
1 step_lse    1152.
2 step_k_fold 1125.
3 ridge       1447.
4 lasso_lse   1236.
5 lasso_min   1179.
6 pcr         1286.
7 pls         1258.

```

- Is there much difference among the test errors resulting from these five approaches?
- Which method appears most accurate?

There does not appear to be a drastic difference in the model accuracy between our approaches. We are able to predict the number of applications with rmse of between 1,150 and 1,300. It looks like the stepwise approach was the most accurate with 12 predictors.

- Look at help for: `pls::validationplot()`, `pls::R2()`, `pls::MSEP()`

```

pls::RMSEP(object = pcr_reg, estimate = "all",
           newdata = model_data$test)

```

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
train	3249	3224	1461	1435	1311	1196	1034
CV	3258	3258	1482	1463	1342	1235	1076
adjCV	3258	3258	1481	1462	1339	1234	1073
test	4401	4404	2155	2096	1973	1742	1286
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps

train	1032	1032	1032	1009	993.4	991.5	982.7
CV	1075	1082	1083	1065	1049.9	1052.6	1047.2
adjCV	1073	1080	1080	1062	1046.8	1049.4	1043.8
test	1276	1270	1270	1221	1193.3	1192.3	1205.9
	14 comps	15 comps	16 comps	17 comps			
train	971.9	964.4	963.9	943			
CV	1041.8	1038.5	1042.4	1023			
adjCV	1037.7	1034.5	1038.1	1018			
test	1181.6	1163.5	1161.0	1172			

```
# this appears to do the same thing that I did earlier manually, but faster. I can pass ne
```

Comparison of Dimension Reduction Methods

We will now try to predict per capita crime rate in the Boston data set in the {MASS} package (`library(MASS)`).

- Leave `rad` as an integer.
- Use `set.seed(1234)` and 50% validation where appropriate.

```
model_data <- split_train_and_test_data(split_pct = .5,
                                         seed = 1234, data = Boston)
```

- Try out at least four of the dimension reduction methods that we explored over the last two weeks, such as the best subset selection, lasso, ridge regression, PCR, and PLS. Discuss results.

Principal Component Analysis Model

```
pcr_model <- pcr(crim ~ ., data = model_data$train, validation = "CV")
summary(pcr_model)
```

```
Data:   X dimension: 253 13
      Y dimension: 253 1
Fit method: svdpc
Number of components considered: 13

VALIDATION: RMSEP
```

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	10.32	8.701	8.735	8.734	8.722	8.660	8.345
adjCV	10.32	8.696	8.726	8.724	8.712	8.647	8.331
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	8.304	8.247	8.233	8.203	8.246	8.248	8.227
adjCV	8.289	8.231	8.217	8.185	8.224	8.225	8.203

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	82.53	96.85	98.99	99.68	99.86	99.92	99.96	99.99
crim	29.95	30.74	30.81	31.12	32.91	38.24	39.15	40.17
	9 comps	10 comps	11 comps	12 comps	13 comps			
X	100.00	100.00	100.00	100.0	100.0			
crim	40.55	41.36	42.07	42.2	42.7			

It appears that 7 components is the optimal number. After that, the increases are extremely marginal. Still, the model is only accounting for ~42% of the variance and the RMSE is pretty high.

RMSE = ~ 8.7

Partial Least Squares Model

```
plsr_model <- plsr(crim ~ ., data = model_data$train, validation = "CV")
summary(plsr_model)
```

Data: X dimension: 253 13

Y dimension: 253 1

Fit method: kernelpls

Number of components considered: 13

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	10.32	8.749	8.763	8.675	8.529	8.362	8.234
adjCV	10.32	8.741	8.752	8.662	8.515	8.348	8.220
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	8.226	8.236	8.217	8.237	8.303	8.318	8.300
adjCV	8.210	8.219	8.200	8.218	8.278	8.291	8.273

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	82.48	96.79	98.18	99.36	99.81	99.92	99.95	99.97
crim	30.22	30.92	33.16	35.68	38.24	40.24	40.79	41.18
	9 comps	10 comps	11 comps	12 comps	13 comps			
X	99.99	100.00	100.00	100.00	100.0			
crim	41.49	41.71	42.18	42.29	42.7			

The pls model performs quite similarly and its pretty hard to tell what the optimal number of components should be. After one component, the changes in variance and RMSE is pretty small.

RMSE = ~8.7

Stepwise Regression

```
y_train <- model_data$train$crim
full_regression <- lm(crim ~ ., data = model_data$train)
step_output <- step(object = full_regression, method = "backward")
```

Start: AIC=1066.32

```
crim ~ zn + indus + chas + nox + rm + age + dis + rad + tax +
      ptratio + black + lstat + medv
```

	Df	Sum of Sq	RSS	AIC
- age	1	0.00	15328	1064.3
- chas	1	22.82	15350	1064.7
- indus	1	23.96	15352	1064.7
- tax	1	29.55	15357	1064.8
- black	1	62.15	15390	1065.3
- ptratio	1	85.09	15413	1065.7
<none>			15328	1066.3
- nox	1	135.92	15464	1066.5
- rm	1	164.66	15492	1067.0
- zn	1	170.49	15498	1067.1
- lstat	1	223.60	15551	1068.0
- dis	1	331.43	15659	1069.7
- medv	1	622.51	15950	1074.4
- rad	1	1320.70	16648	1085.2

Step: AIC=1064.32


```
crim ~ zn + indus + chas + nox + rm + dis + rad + tax + ptratio +
      black + lstat + medv
```

	Df	Sum of Sq	RSS	AIC
- chas	1	22.86	15350	1062.7
- indus	1	23.96	15352	1062.7
- tax	1	29.55	15357	1062.8
- black	1	62.21	15390	1063.3
- ptratio	1	85.37	15413	1063.7
<none>			15328	1064.3
- nox	1	146.70	15474	1064.7
- rm	1	174.32	15502	1065.2
- zn	1	177.64	15505	1065.2
- lstat	1	240.94	15568	1066.3
- dis	1	357.01	15685	1068.1
- medv	1	625.62	15953	1072.4
- rad	1	1328.81	16656	1083.3

Step: AIC=1062.69

```
crim ~ zn + indus + nox + rm + dis + rad + tax + ptratio + black +
      lstat + medv
```

	Df	Sum of Sq	RSS	AIC
- tax	1	25.43	15376	1061.1
- indus	1	27.71	15378	1061.2
- black	1	65.67	15416	1061.8
- ptratio	1	84.66	15435	1062.1
<none>			15350	1062.7
- nox	1	164.37	15515	1063.4
- zn	1	177.55	15528	1063.6
- rm	1	178.03	15528	1063.6
- lstat	1	241.90	15592	1064.7
- dis	1	362.41	15713	1066.6
- medv	1	663.50	16014	1071.4
- rad	1	1312.14	16662	1081.4

Step: AIC=1061.11

```
crim ~ zn + indus + nox + rm + dis + rad + ptratio + black +
      lstat + medv
```

	Df	Sum of Sq	RSS	AIC
- black	1	64.59	15440	1060.2
- indus	1	91.95	15468	1060.6

- ptratio	1	93.80	15470	1060.7
<none>			15376	1061.1
- zn	1	156.90	15533	1061.7
- nox	1	169.01	15545	1061.9
- rm	1	171.53	15547	1061.9
- lstat	1	244.71	15620	1063.1
- dis	1	363.02	15739	1065.0
- medv	1	639.24	16015	1069.4
- rad	1	2698.70	18074	1100.0

Step: AIC=1060.17

crim ~ zn + indus + nox + rm + dis + rad + ptratio + lstat +
medv

	Df	Sum of Sq	RSS	AIC
- indus	1	87.57	15528	1059.6
- ptratio	1	110.00	15550	1060.0
<none>			15440	1060.2
- zn	1	162.05	15602	1060.8
- nox	1	179.36	15620	1061.1
- rm	1	219.95	15660	1061.8
- lstat	1	282.44	15723	1062.8
- dis	1	379.04	15820	1064.3
- medv	1	741.43	16182	1070.0
- rad	1	2984.79	18425	1102.9

Step: AIC=1059.6

crim ~ zn + nox + rm + dis + rad + ptratio + lstat + medv

	Df	Sum of Sq	RSS	AIC
<none>			15528	1059.6
- ptratio	1	132.06	15660	1059.8
- zn	1	176.56	15705	1060.5
- rm	1	251.19	15779	1061.7
- lstat	1	273.81	15802	1062.0
- nox	1	274.82	15803	1062.0
- dis	1	308.96	15837	1062.6
- medv	1	695.31	16223	1068.7
- rad	1	2916.19	18444	1101.2

```

step_summary <- summary(step_output)

step_best_model <- lm(formula = crim ~ zn + nox + rm +
                      dis + rad + ptratio +
                      lstat + medv, data = model_data$train)

step_predictions <- predict(object = step_best_model, newdata = model_data$train)

sqrt(calc_mse(y_actual = y_train, y_predicted = step_predictions))

```

```
[1] 7.834251
```

Stepwise regression selection selected 7 variables from the original 12 and resulted with rmse of 7.83.

Ridge Regression

```

x_train <- model.matrix(lm(formula = crim ~ .,
                          data = model_data$train))[, -1]

r_train <- cv.glmnet(x = x_train, y = y_train, alpha = 0)

ridge_predictions <- predict(r_train, newx = x_train) %>% as.vector()
ridge_predictions

```

```

[1] 1.612530 3.283893 5.885659 3.807628 5.172257 3.927413 5.889116 4.456629
[9] 5.632322 5.989434 6.057029 2.824525 4.544568 3.335639 3.103409 3.144530
[17] 3.692734 5.331243 4.800348 3.385360 5.617881 3.174858 3.730548 2.699924
[25] 4.010553 5.075539 2.573706 4.032486 5.753436 2.423660 3.429051 4.158854
[33] 6.069259 2.625611 3.561721 6.143913 5.654299 5.469102 2.753608 4.151851
[41] 4.543778 3.312781 2.392795 4.080233 3.428088 3.501909 6.448969 3.504647
[49] 2.626781 2.826792 3.998475 5.539902 3.121006 2.862221 5.223336 4.928210
[57] 3.051871 5.728671 5.458974 3.536729 3.493848 4.034943 3.916215 4.669889
[65] 2.569136 5.181098 4.284687 3.397671 2.940994 3.161263 3.972334 2.660683
[73] 3.135902 3.808443 5.172708 2.478554 2.470502 3.634692 3.725109 3.406174
[81] 4.822377 4.662451 5.857774 3.022307 3.563799 5.873913 3.601747 3.941436
[89] 5.717717 2.811296 4.250093 3.124699 3.517978 2.687095 6.255103 5.794497
[97] 4.661772 5.583352 5.222162 3.335123 3.675303 6.200523 3.925205 4.731010

```

```
[105] 4.173565 3.684115 3.660640 4.307700 5.754074 3.737641 4.875767 3.720710
[113] 5.535622 3.187702 4.251081 4.390558 3.860719 3.492540 5.702490 3.083529
[121] 4.814622 2.952592 3.483246 3.872219 2.761304 3.268598 5.165511 2.738989
[129] 5.569939 4.644664 3.256783 5.739562 3.204553 6.102026 5.185986 3.921102
[137] 6.382558 3.524596 3.140672 3.311948 5.906258 5.415257 6.005334 3.236285
[145] 3.113515 5.289821 5.127521 6.737446 2.946142 2.879296 4.256979 4.756471
[153] 6.294249 4.227496 3.868277 2.988178 3.484438 5.770035 5.274886 2.875234
[161] 1.960236 3.150972 5.415163 2.926357 4.849388 2.467219 5.475836 5.037559
[169] 3.841829 6.034724 5.470242 6.039459 3.611647 3.144789 5.133556 3.540403
[177] 2.704194 5.344277 3.003421 5.455231 4.288729 2.710407 5.364268 4.305857
[185] 3.012002 4.361344 2.930048 2.885458 4.669527 2.895473 3.405247 4.448731
[193] 3.414483 5.575409 3.432226 3.397185 5.961657 2.255470 6.081084 3.062482
[201] 3.775953 4.205282 3.607564 3.255160 5.554167 3.115011 5.847252 3.205094
[209] 3.087763 5.146055 2.501293 2.083660 3.182924 3.149880 2.568591 3.522708
[217] 2.410190 3.457947 2.017172 2.265563 3.633992 3.547817 1.954263 4.116003
[225] 4.037627 4.104865 5.003165 5.566102 3.277276 5.152197 3.324057 5.541940
[233] 3.673378 6.057449 3.874799 2.350740 3.632658 3.390823 2.656352 5.252871
[241] 4.129647 2.326829 3.166769 3.328426 5.568014 5.473942 3.477075 5.519424
[249] 4.106352 3.355919 4.239522 2.832617 3.198019
```

```
sqrt(calc_mse(y_actual = y_train, y_predicted = ridge_predictions))
```

```
[1] 9.683361
```

The ridge regression doesn't remove any variables and the rmse is 9.46.

Overall, there was not a huge difference in performance. The stepwise selected model performed the best, with an rmse of ~ 7.9. The PCR and PLSR model perform similarly and the ridge regression did the worst.

Create Plots

- a. For all possible values of K, compare
 - regression models that are based on the best subset of p -variables.

```
n_predictors <- ncol(model_data$train) - 1

# empty data frame to store the results.
results <- data.frame(matrix(NA, ncol = 3, nrow = n_predictors))
colnames(results) <- c("Variables", "adjusted_r2", "rmse")
```

```

for(k in 1:n_predictors) {
  #browser()
  predictor_combinations <- regsubsets(crim ~ .,
                                       data = model_data$train,
                                       nvmax = k, intercept = FALSE)

  # Fit models for each combination and obtain adjusted R-squared and MSE
  models <- summary(predictor_combinations)
  adj_r_squared_values <- models$rsq

  # Find the index of the model with the highest adjusted R-squared
  best_model_index <- which.max(adj_r_squared_values)

  # Extract the names of selected predictors
  selected_predictors <- names(coef(predictor_combinations, id = best_model_index))

  # Fit a linear model using the selected predictors
  lm_model <- lm(crim ~ .,
                data = model_data$train[,c("crim", selected_predictors)])

  # Calculate MSE
  y_predicted <- predict(lm_model, newdata = model_data$train)
  rmse_value <- sqrt(mean((model_data$train$crim - y_predicted)^2))

  # Store results for the current k
  results[k, 1] <- length(names(coef(predictor_combinations, id = best_model_index)))
  results[k, 2] <- adj_r_squared_values[best_model_index]
  results[k, 3] <- rmse_value
}

step_results <- results
step_results

```

	Variables	adjusted_r2	rmse
1	1	0.4090749	8.318242
2	2	0.4545558	8.121705
3	3	0.4709485	8.044823
4	4	0.4744398	8.018398
5	5	0.4811558	7.956229

6	6	0.4870767	7.921192
7	7	0.4938450	7.867494
8	8	0.4977103	7.838929
9	9	0.4999329	7.819515
10	10	0.5010587	7.795773
11	11	0.5018515	7.789602
12	12	0.5025528	7.783521
13	13	0.5025659	7.783521

- PCR based on first principal components;
- The `pls::R2()` and `pls::MSEP()` functions calculate this automatically, so I will just extract those values as a vector.

```

pcr_model <- pcr(crim ~ . -1, data = model_data$train,
                 validation = "CV", scaled = TRUE)

pcr_r_squared <- pls::R2(pcr_model, estimate = "test",
                        newdata = model_data$test)$val %>%
data.frame() %>%
  pivot_longer(cols = everything(),
               names_to = "component",
               values_to = "r_squared") %>%
  slice(-1)

pcr_rmse <- pls::RMSEP(pcr_model, estimate = "test",
                      newdata = model_data$test)$val %>%
data.frame() %>%
  pivot_longer(cols = everything(),
               names_to = "component",
               values_to = "rmse") %>%
  slice(-1) %>%
  dplyr::select(rmse)

pcr_results <- bind_cols(pcr_r_squared, pcr_rmse)

```

- PLS based on first PLS components.

```

plsr_model <- plsr(crim ~ ., data = model_data$train,
                  validation = "CV", scaled = TRUE)

plsr_r_squared <- pls::R2(plsr_model, estimate = "test",
                        newdata = model_data$test)$val %>%

```

```

data.frame() %>%
  pivot_longer(cols = everything(),
               names_to = "component",
               values_to = "r_squared") %>%
  slice(-1)

plsr_rmse <- pls::RMSEP(plsr_model, estimate = "test",
                      newdata = model_data$test)$val %>%
data.frame() %>%
  pivot_longer(cols = everything(),
               names_to = "component",
               values_to = "rmse") %>%
  slice(-1) %>%
  dplyr::select(rmse)

plsr_results <- bind_cols(plsr_r_squared, plsr_rmse)

```

For each , compare these methods in terms of the **explained proportion of the total variation of crime rate per capita** (adjusted) and in terms of the **prediction mean-squared error**.


- Hint: Look at the regsubsets output object list elements. The `$adjR2` will be useful. You may find it helpful to use the `$which` matrix in a for-loop to get the cross-validated MSE.
- c. Create one plot comparing adjusted- R^2 for different values of k for subset regression, PCR, and PLS and interpret the plot.

```

r_squared_data <- cbind(vars = step_results$Variables,
                      step_r2 = step_results$adjusted_r2,
                      pcr_r2 = pcr_results$r_squared,
                      plsr_r2 = plsr_results$r_squared) %>%
  as.data.frame()

ggplot(data = r_squared_data, aes(x = vars, y = step_r2)) +
  geom_line(color = "red") +
  geom_line(mapping = aes(y = pcr_r2), color = "blue") +
  geom_line(mapping = aes(y = plsr_r2), color = "purple") +
  labs(x = "Number of Variables", y = "Adjusted R-Squared")

```




Based on the r-squared values, the stepwise model selection performed best with a max adjusted r-squared value of roughly 50% when it used all of the variables. Interestingly, the partial least squared model was outperforming the PCR model until the number of dimensions exceeded 8 or so. I am guessing this is an issue with scaling the variables.

- c. Create one plot comparing Prediction MSE for different values of k for subset regression, PCR, and PLS and interpret the plot.

```
rmse_data <- cbind(vars = step_results$Variables,
                  step_rmse = step_results$rmse,
                  pcr_rmse = pcr_results$rmse,
                  plsr_rmse = plsr_results$rmse) %>%
  as.data.frame()

ggplot(data = rmse_data, aes(x = vars, y = step_rmse)) +
  geom_line(color = "red") +
  geom_line(mapping = aes(y = pcr_rmse), color = "blue") +
  geom_line(mapping = aes(y = plsr_rmse), color = "purple") +
  labs(x = "Number of Variables", y = "RMSE")
```

- e. Propose a model (or set of models) that seem to perform well on this data set, and justify your answer. Make sure you are evaluating model performance on basis of cross-validation, as opposed to using training error.
- f. Does your chosen model involve all of the features in the data set? Why or why not?

I am not entirely sure that I set this up correctly. I don't fully understand estimate argument in the PCR/PLS function. I have changed it several times and get widely different results. It decided to make the estimate "test" which given that the model was created with the training makes sense to me.

I would select the PCR models with somewhere between 2-3 components. As we can see from the plot, the RMSE is pretty stable for with 2 components vs 10; therefore, its not worth the degrees-of-freedom price we pay to get an extra .001 in precision in the rmse.